

# Reinforcement Learning Maze Runner: When RL Meets Real Human Action

Name : Sriman Soma (S20230030409)



# RECAP (Overview of the Project)



## Environment

- $10 \times 10$  maze
- Start → Goal navigation
- Walls + Traps (Potholes)
- 100 discrete states, 4 actions

## Rewards

- **+100** Goal
- **-50** Pothole (Added this to make differ from simple grid world)
- **-10** Wall / invalid move
- **-0.1** Normal step cost

## Algorithms Implemented so far:

- Value Iteration — as Assignment 2
- Monte Carlo Control — as Assignment 3
- SARSA — as Assignment 3
- Q-Learning (Final chosen) — as Assignment 3
- $Q(\lambda)$  — as Assignment 3
- Linear Function Approximation (explored additionally)



# Reinforcement Learning

## Learning by Experience

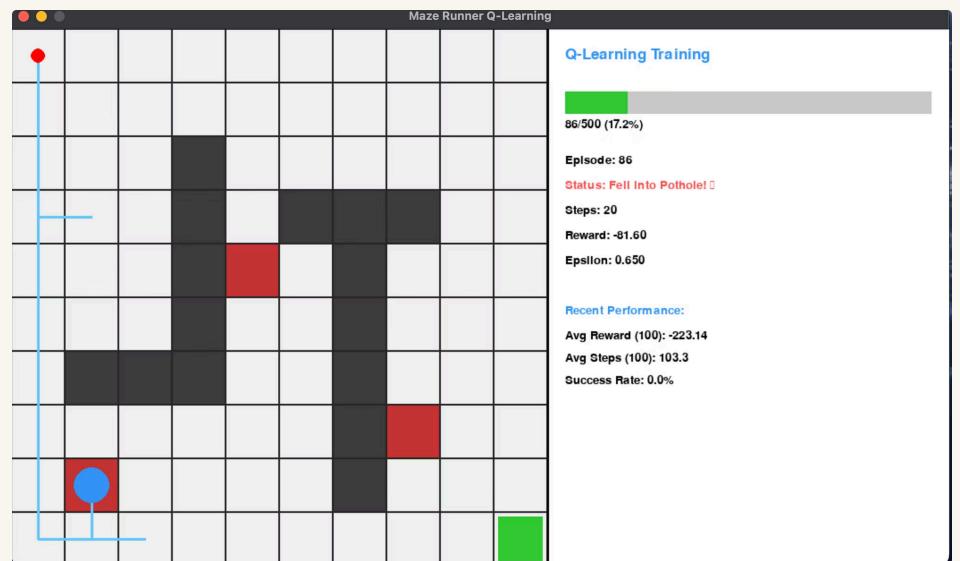


### From Zero to Intelligence

- Q-Learning agent starts with **zero knowledge (Q-table = 0)**.
- Learns optimal policy **purely from experience – no maze map**.
- **Bellman update:**  
$$Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma \max Q(s',a') - Q(s,a)]$$
  
( $\alpha = 0.1, \gamma = 0.99$  → balances immediate vs future rewards)
- **Exploration vs Exploitation**
  - $\epsilon$ -greedy policy with decay:  $\epsilon = 1.0 \rightarrow 0.01$  ( $\times 0.995$  per episode)
  - **Early:** full exploration (random moves)
  - **Mid:** mix of explore + exploit
  - **Late:** pure exploitation (optimal path)
- **Continuous Improvement**
  - **Ep 1–100:** Random, high variance (~200 + steps)
  - **Ep 100–300:** Q-values stabilize, 60–70 % success
  - **Ep 300–500:** Optimal policy, 18–20 steps, > 95 % success

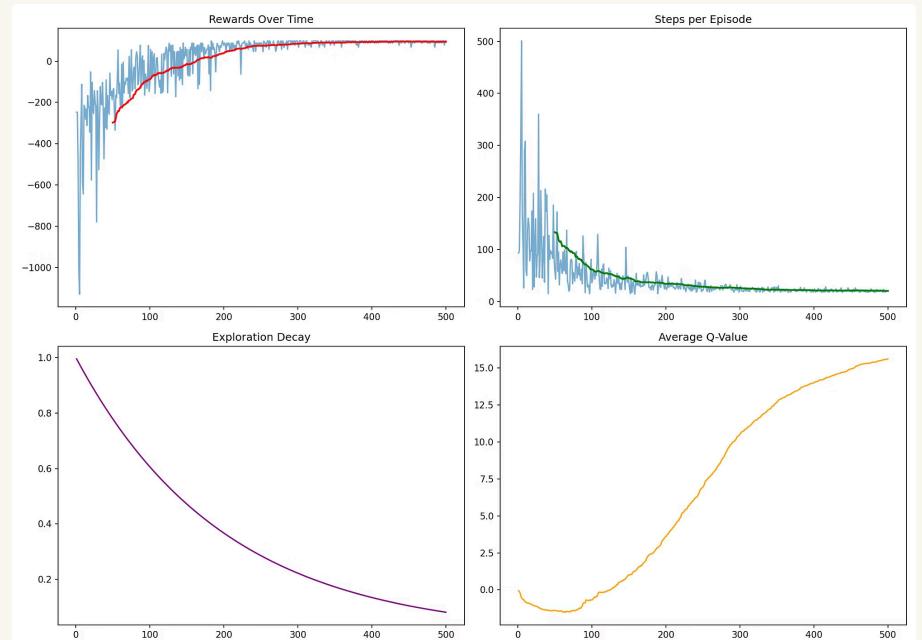
# Unique features I implemented in this project:

- Pothole termination logic
- Strong negative reward  $-50$  for traps
- Immediate episode termination on entering trap
- Agent restarts fresh next episode
- Red trap visualization in GUI
- Live status indicator (“Fell into Pothole!”)



## Why this improves learning

- Avoids wasting steps inside dead zones
- Encourages safe path discovery
- Speeds up Q-value convergence
- Creates clearer patterns in training video



# Why Q-Learning?

- Type: **Model-Free** | **Off-Policy** | **Value-Based**

## Key Strengths

- **Model-Free:** Learns directly from experience — no need for transition model  $P(s' | s, a)P(s' | s,a)$ 
  - **Off-Policy:** Learns optimal (greedy) policy while exploring ( $\epsilon$ -greedy)
- **Fast TD Updates:** Updates after every step (bootstrapping) → faster & low-variance learning

## Perfect for the Maze Environment

Simple Tabular Setup:  $10 \times 10$  grid  $\times 4$  actions → only 400 Q-values

- **Sample Efficient (TD(0)):**  
One-step updates reuse learned values → ideal for short episodes (~500 steps)
- **Theoretical Convergence:**  
Finite state-action space  
 $\epsilon$ -greedy exploration

# Why Not Other Algorithms?

## 🚫 Dynamic Programming

- Needs full  $P(s'|s,a)$  and  $R(s,a)$  model — **not available** in our maze
- Suited for known environments, not learning tasks

## 🚫 Monte Carlo

- Must complete entire episode before update
- High variance, inefficient for long (500-step) episodes

## 🚫 SARSA

- On-policy → learns from behavior policy
- Slower convergence, overly cautious in exploration

## 🚫 Deep Q-Networks (DQN)

- Adds unnecessary neural-network complexity
- Risk of instability
- Tabular Q-Learning already gives exact solution

## 🚫 Function Approximation Makes Sense :

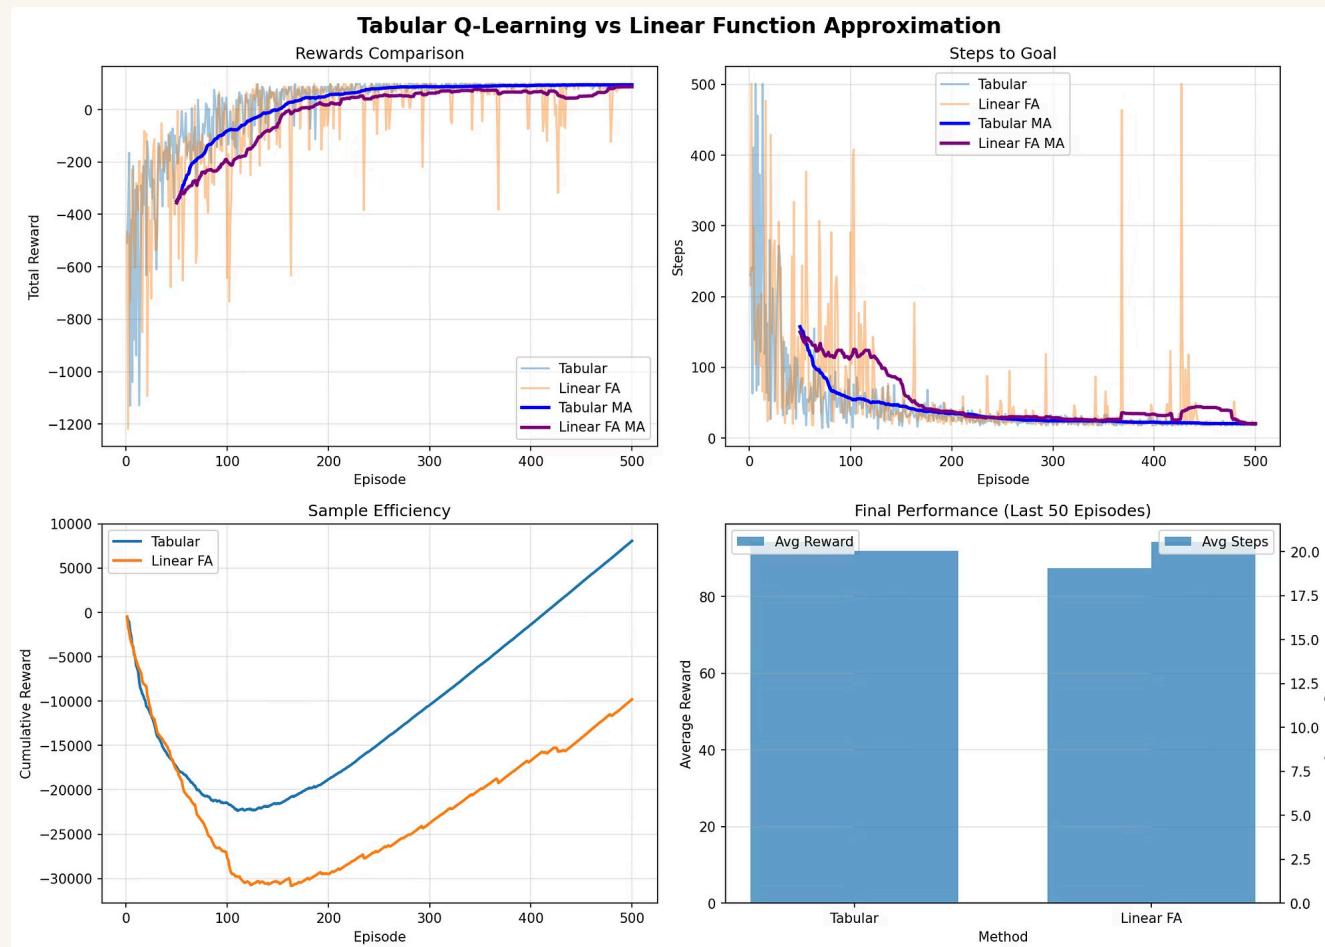
- Large/continuous state spaces: e.g.,  $1000 \times 1000$  maze, pixel inputs, robot joint angles
- Generalization needed: Learning from similar states
- Memory constraints: Can't store millions of Q-values

## 🚫 Linear Function Approximation Makes Sense :

Small & Discrete State Space → Tabular Is Exact

- Only 100 states, 4 actions → 400 Q-values total
- Tabular Q-learning gives exact values
- Linear FA introduces approximation error

# Why the Linear Function approximation is not good in this project:



- Linear FA gave very unstable results and the rewards went extremely low.
- The number of steps kept jumping a lot, meaning it did not learn properly.
- It learned very slowly and needed many more episodes to improve.
- Even after full training, Linear FA performed worse than normal tabular Q-learning.

TRAINING VIDEO LINK for Q-learning:

[https://drive.google.com/file/d/1DoycAMmUd5eobK6NZCEm2E-uVsWJOx4I/view?  
usp=sharing](https://drive.google.com/file/d/1DoycAMmUd5eobK6NZCEm2E-uVsWJOx4I/view?usp=sharing)